

Lab08_Mini_Assignment

Matthew

In this class we will explore a complete analysis using the unsupervised learning techniques

```
wisc.df <- read.csv("WisconsinCancer.csv", row.names=1)
head(wisc.df)
```

##	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
## 842302	M	17.99	10.38	122.80	1001.0
## 842517	M	20.57	17.77	132.90	1326.0
## 84300903	M	19.69	21.25	130.00	1203.0
## 84348301	M	11.42	20.38	77.58	386.1
## 84358402	M	20.29	14.34	135.10	1297.0
## 843786	M	12.45	15.70	82.57	477.1
##	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
## 842302	0.11840	0.27760	0.3001	0.14710	
## 842517	0.08474	0.07864	0.0869	0.07017	
## 84300903	0.10960	0.15990	0.1974	0.12790	
## 84348301	0.14250	0.28390	0.2414	0.10520	
## 84358402	0.10030	0.13280	0.1980	0.10430	
## 843786	0.12780	0.17000	0.1578	0.08089	
##	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
## 842302	0.2419	0.07871	1.0950	0.9053	8.589
## 842517	0.1812	0.05667	0.5435	0.7339	3.398
## 84300903	0.2069	0.05999	0.7456	0.7869	4.585
## 84348301	0.2597	0.09744	0.4956	1.1560	3.445
## 84358402	0.1809	0.05883	0.7572	0.7813	5.438
## 843786	0.2087	0.07613	0.3345	0.8902	2.217
##	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
## 842302	153.40	0.006399	0.04904	0.05373	0.01587
## 842517	74.08	0.005225	0.01308	0.01860	0.01340
## 84300903	94.03	0.006150	0.04006	0.03832	0.02058
## 84348301	27.23	0.009110	0.07458	0.05661	0.01867
## 84358402	94.44	0.011490	0.02461	0.05688	0.01885
## 843786	27.19	0.007510	0.03345	0.03672	0.01137
##	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
## 842302	0.03003	0.006193	25.38	17.33	
## 842517	0.01389	0.003532	24.99	23.41	
## 84300903	0.02250	0.004571	23.57	25.53	
## 84348301	0.05963	0.009208	14.91	26.50	
## 84358402	0.01756	0.005115	22.54	16.67	
## 843786	0.02165	0.005082	15.47	23.75	
##	perimeter_worst	area_worst	smoothness_worst	compactness_worst	
## 842302	184.60	2019.0	0.1622	0.6656	
## 842517	158.80	1956.0	0.1238	0.1866	
## 84300903	152.50	1709.0	0.1444	0.4245	

```
## 84348301      98.87      567.7      0.2098      0.8663
## 84358402     152.20     1575.0      0.1374      0.2050
## 843786       103.40      741.6      0.1791      0.5249
##      concavity_worst concave.points_worst symmetry_worst
## 842302      0.7119      0.2654      0.4601
## 842517      0.2416      0.1860      0.2750
## 84300903     0.4504      0.2430      0.3613
## 84348301     0.6869      0.2575      0.6638
## 84358402     0.4000      0.1625      0.2364
## 843786      0.5355      0.1741      0.3985
##      fractal_dimension_worst
## 842302      0.11890
## 842517      0.08902
## 84300903     0.08758
## 84348301     0.17300
## 84358402     0.07678
## 843786      0.12440
```

Remove the diagnosis column and keep it in a separate vector for later

```
diagnosis <- as.factor(wisc.df[,1])
wisc.data <- wisc.df[,-1]
head(wisc.data)
```

```
##      radius_mean texture_mean perimeter_mean area_mean smoothness_mean
## 842302      17.99      10.38      122.80      1001.0      0.11840
## 842517      20.57      17.77      132.90      1326.0      0.08474
## 84300903     19.69      21.25      130.00      1203.0      0.10960
## 84348301     11.42      20.38       77.58       386.1      0.14250
## 84358402     20.29      14.34      135.10      1297.0      0.10030
## 843786      12.45      15.70       82.57       477.1      0.12780
##      compactness_mean concavity_mean concave.points_mean symmetry_mean
## 842302      0.27760      0.3001      0.14710      0.2419
## 842517      0.07864      0.0869      0.07017      0.1812
## 84300903     0.15990      0.1974      0.12790      0.2069
## 84348301     0.28390      0.2414      0.10520      0.2597
## 84358402     0.13280      0.1980      0.10430      0.1809
## 843786      0.17000      0.1578      0.08089      0.2087
##      fractal_dimension_mean radius_se texture_se perimeter_se area_se
## 842302      0.07871      1.0950      0.9053      8.589 153.40
## 842517      0.05667      0.5435      0.7339      3.398 74.08
## 84300903     0.05999      0.7456      0.7869      4.585 94.03
## 84348301     0.09744      0.4956      1.1560      3.445 27.23
## 84358402     0.05883      0.7572      0.7813      5.438 94.44
## 843786      0.07613      0.3345      0.8902      2.217 27.19
##      smoothness_se compactness_se concavity_se concave.points_se
## 842302      0.006399      0.04904      0.05373      0.01587
## 842517      0.005225      0.01308      0.01860      0.01340
## 84300903     0.006150      0.04006      0.03832      0.02058
## 84348301     0.009110      0.07458      0.05661      0.01867
## 84358402     0.011490      0.02461      0.05688      0.01885
## 843786      0.007510      0.03345      0.03672      0.01137
##      symmetry_se fractal_dimension_se radius_worst texture_worst
```

```
## 842302      0.03003      0.006193      25.38      17.33
## 842517      0.01389      0.003532      24.99      23.41
## 84300903    0.02250      0.004571      23.57      25.53
## 84348301    0.05963      0.009208      14.91      26.50
## 84358402    0.01756      0.005115      22.54      16.67
## 843786      0.02165      0.005082      15.47      23.75
##      perimeter_worst area_worst smoothness_worst compactness_worst
## 842302      184.60      2019.0      0.1622      0.6656
## 842517      158.80      1956.0      0.1238      0.1866
## 84300903    152.50      1709.0      0.1444      0.4245
## 84348301      98.87      567.7      0.2098      0.8663
## 84358402    152.20      1575.0      0.1374      0.2050
## 843786      103.40      741.6      0.1791      0.5249
##      concavity_worst concave.points_worst symmetry_worst
## 842302      0.7119      0.2654      0.4601
## 842517      0.2416      0.1860      0.2750
## 84300903    0.4504      0.2430      0.3613
## 84348301    0.6869      0.2575      0.6638
## 84358402    0.4000      0.1625      0.2364
## 843786      0.5355      0.1741      0.3985
##      fractal_dimension_worst
## 842302      0.11890
## 842517      0.08902
## 84300903    0.08758
## 84348301    0.17300
## 84358402    0.07678
## 843786      0.12440
```

Q1. How many observations are in this dataset?

```
nrow(wisc.data)
```

```
## [1] 569
```

There are 569 observations

Q2. How many of the observations have a malignant diagnosis?

```
table(wisc.df$diagnosis)
```

```
##
##      B      M
## 357 212
```

212 Malignant Diagnosis

Q3. How many variables/features in the data are suffixed with `__mean`?

First find the column names

```
colnames(wisc.data)
```

```
## [1] "radius_mean"          "texture_mean"
## [3] "perimeter_mean"       "area_mean"
## [5] "smoothness_mean"      "compactness_mean"
## [7] "concavity_mean"       "concave.points_mean"
## [9] "symmetry_mean"        "fractal_dimension_mean"
## [11] "radius_se"           "texture_se"
## [13] "perimeter_se"         "area_se"
## [15] "smoothness_se"        "compactness_se"
## [17] "concavity_se"         "concave.points_se"
## [19] "symmetry_se"          "fractal_dimension_se"
## [21] "radius_worst"         "texture_worst"
## [23] "perimeter_worst"      "area_worst"
## [25] "smoothness_worst"     "compactness_worst"
## [27] "concavity_worst"      "concave.points_worst"
## [29] "symmetry_worst"       "fractal_dimension_worst"
```

Next I need to search within the column names for “`__mean`” pattern. The ‘`grep()`’ function might.

```
grep("__mean",colnames(wisc.data))
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
inds <- grep("__mean",colnames(wisc.data))
length(inds)
```

```
## [1] 10
```

There are 10 variables with suffix `__mean`

Q. How many dimesnions are in this dataset?

```
ncol(wisc.data)
```

```
## [1] 30
```

Principal Component Analysis

First do we need to scale the data before PCA or not

```
round(apply(wisc.data,2,sd),3)
```

```
##           radius_mean      texture_mean      perimeter_mean
##           3.524         4.301         24.299
##           area_mean      smoothness_mean      compactness_mean
##           351.914         0.014         0.053
##           concavity_mean      concave.points_mean      symmetry_mean
##           0.080         0.039         0.027
## fractal_dimension_mean      radius_se      texture_se
##           0.007         0.277         0.552
##           perimeter_se      area_se      smoothness_se
##           2.022         45.491         0.003
##           compactness_se      concavity_se      concave.points_se
##           0.018         0.030         0.006
##           symmetry_se      fractal_dimension_se      radius_worst
##           0.008         0.003         4.833
##           texture_worst      perimeter_worst      area_worst
##           6.146         33.603         569.357
##           smoothness_worst      compactness_worst      concavity_worst
##           0.023         0.157         0.209
##           concave.points_worst      symmetry_worst      fractal_dimension_worst
##           0.066         0.062         0.018
```

Looks like we need to scale.

```
wisc.pr <- prcomp(wisc.data,scale=TRUE)
summary(wisc.pr)
```

```
## Importance of components:
##           PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  3.6444 2.3857 1.67867 1.40735 1.28403 1.09880 0.82172
## Proportion of Variance 0.4427 0.1897 0.09393 0.06602 0.05496 0.04025 0.02251
## Cumulative Proportion 0.4427 0.6324 0.72636 0.79239 0.84734 0.88759 0.91010
##           PC8      PC9      PC10      PC11      PC12      PC13      PC14
## Standard deviation  0.69037 0.6457 0.59219 0.5421 0.51104 0.49128 0.39624
## Proportion of Variance 0.01589 0.0139 0.01169 0.0098 0.00871 0.00805 0.00523
## Cumulative Proportion 0.92598 0.9399 0.95157 0.9614 0.97007 0.97812 0.98335
##           PC15      PC16      PC17      PC18      PC19      PC20      PC21
## Standard deviation  0.30681 0.28260 0.24372 0.22939 0.22244 0.17652 0.1731
## Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
## Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
##           PC22      PC23      PC24      PC25      PC26      PC27      PC28
## Standard deviation  0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
## Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
## Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
##           PC29      PC30
## Standard deviation  0.02736 0.01153
## Proportion of Variance 0.00002 0.00000
## Cumulative Proportion 1.00000 1.00000
```

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

44.27%

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

3 PCs capture 72%

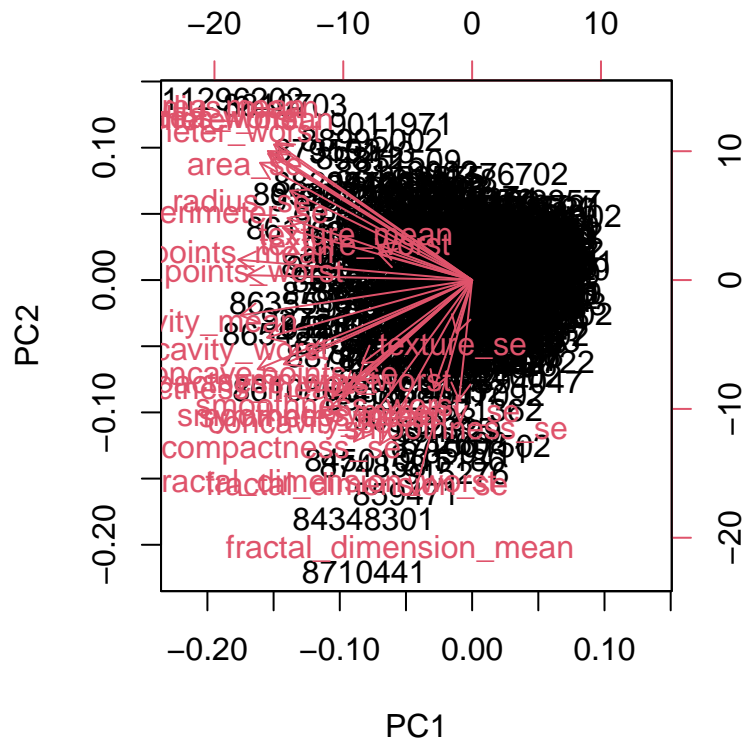
Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

7 PCs capture 91%

PC Plot

Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?


```
biplot(wisc.pr)
```



It is very hard to understand because there is a lot of data overlapping each other as well as the labels for this graph. This compares the PC1 and PC2.


We need to make our plot of PC1 vs PC2 (aka score plot, PC-plot, etc.)/ The main result of PCA...

```
plot(wisc.pr$x[,1],wisc.pr$x[,2],col=diagnosis)
```




```
library(ggplot2)
pc <- as.data.frame(wisc.pr$x)
pc$diagnosis <- diagnosis

ggplot(pc)+
  aes(PC1,PC2,col=diagnosis)+
  geom_point()
```

Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
plot(wisc.pr$x[,1],wisc.pr$x[,3],col=diagnosis)
```



The samples for the benign and malignant seem to be more mixed in with each other.

Calculate variance of each component

```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

```
## [1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```


Variance explained by each principal component: pve

```
pve <- pr.var / sum(pr.var)
head(pve)
```


```
## [1] 0.44272026 0.18971182 0.09393163 0.06602135 0.05495768 0.04024522
```

Plot variance explained for each principal component

```
plot(pve, xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     ylim = c(0, 1), type = "o")
```



```
# Alternative scree plot of the same data, note data driven y-axis
barplot(pve, ylab = "Precent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```



#Examine the PC loadings How much do the original variables contribute to the new PCs that we have calculated? To get at this data we can look at the 'rotation' component of the returned PCA object

```
head(wisc.pr$rotation[,1:3])
```

```
##              PC1              PC2              PC3
## radius_mean   -0.2189024  0.23385713 -0.008531243
## texture_mean  -0.1037246  0.05970609  0.064549903
## perimeter_mean -0.2275373  0.21518136 -0.009314220
## area_mean     -0.2209950  0.23107671  0.028699526
## smoothness_mean -0.1425897 -0.18611302 -0.104291904
## compactness_mean -0.2392854 -0.15189161 -0.074091571
```

Focus in on PC1

```
head(wisc.pr$rotation[,1])
```

```
##      radius_mean      texture_mean      perimeter_mean      area_mean
##      -0.2189024      -0.1037246      -0.2275373      -0.2209950
## smoothness_mean compactness_mean
##      -0.1425897      -0.2392854
```

#Q9. For the first principal component, what is the component of the loading vector (i.e. wisc.pr\$rotation[,1]) for the feature concave.points_mean?

```
wisc.pr$rotation["concave.points_mean",1]
```


```
## [1] -0.2608538
```

This is a complicated mix of variables that go together to make up PC1 - ie there are many of the original variables that together contribute highly to PC1.

#Q10. What is the minimum number of principal components required to explain 80% of the variance of the data? Atleast 3

```
loadings <- as.data.frame(wisc.pr$rotation)
```

```
ggplot(loadings)+
  aes(PC1,rownames(loadings))+
  geom_col()
```




Hierarchical Clustering

```
# Scale the wisc.data data using the "scale()" function
data.scaled <- scale(wisc.data)
```

First we scale the data, then distance matrix, then hclust

```
wisc.hclust <- hclust(dist(scale(data.scaled)))
```

```
plot(wisc.hclust)
```



```
dist(scale(data.scaled))
hclust (*, "complete")
```

Cut this tree to yield our cluster membership vector with 'cutree()' function.

```
grps <- cutree(wisc.hclust, h=19)
table(grps)
```


```
## grps
##  1  2  3  4
## 177  7 383  2
```

```
table(grps,diagnosis)
```

```
##      diagnosis
## grps    B    M
##   1  12 165
##   2   2   5
##   3 343  40
##   4   0   2
```

Q11. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters?

```
plot(wisc.hclust)
abline(wisc.hclust, col="red", lty=2)
```




```
dist(scale(data.scaled))
hclust (*, "complete")
```

At a height of 19

Combine methods: PCA and HCLUST

My PCA results were interesting as they showed a separation of M and B samples along PC1.

```
plot(wisc.pr$x[,1], wisc.pr$x[,2], col=diagnosis)
```



Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?

Increasing the height would mean that the cluster vs diagnoses match would become better.


I want to cluster my PCA results - that is use 'wisc.pr\$x' as input to 'hclust()'

Q13. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

Try clustering in 3 PCs, that is PC1, PC2, and PC3 as input

```
d <- dist(wisc.pr$x[,1:3])
wisc.pr.clust <- hclust(d,method="ward.D2")
plot(wisc.pr.clust)
```


Cluster Dendrogram



```
d  
hclust(*, "ward.D2")
```

I like ward.D2 since that data seems more spread out and clear/organized so it is easier to see where the lines run to on the graph.


Q13. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

Let's cut this tree into two groups/clusters

```
grps <- cutree(wisc.pr.clust,k=2)  
table(grps)
```

```
## grps  
## 1 2  
## 203 366
```

```
plot(wisc.pr$x[,1], wisc.pr$x[,2], col=grps)
```



How well do the two clusters separate the M and B diagnoses

```
table(grps,diagnosis)
```

```
##      diagnosis
## grps    B    M
##    1  24 179
##    2 333  33
```

Q15. How well does the newly created model with four clusters separate out the two diagnoses?

```
(179+333)/nrow(wisc.data)
```

```
## [1] 0.8998243
```

Around 90%